

Winning the Mid-Market: A Blueprint for Sundial

May 2026 | Paras Doshi

Core thesis: Sundial wins the mid-market by selling leverage, not intelligence. The wedge is not replacing strong data teams; it is removing the noise that keeps those teams from doing high-value work.

Executive summary

Companies in this segment (\$100M-\$1B revenue, 10-30 person data teams) already have strong analytics talent. They use Snowflake, dbt, and Airflow. They run sophisticated ML models. They do not need a tool that thinks for them. They need a tool that removes the noise so their expensive humans can focus on work that moves the business.

Picking the mid-market is not just a GTM decision. It is a product strategy decision. It tells Sundial what to build next, what to defer, and what to stop building entirely. That clarity is the real value of choosing a segment.

This document covers four things:

1. The two non-negotiables to land the Opendoor contract
2. Why the mid-market is the right primary target, and how it simplifies the product roadmap
3. A competitive market map with fact-checked vendor placements
4. The sequencing playbook based on which persona Sundial pursues

1. The Opendoor contract: two non-negotiables

Opendoor runs a data organization with pods covering acquisition, conversion, resale, AI analytics, and data engineering. The team builds direct mail targeting models, seller funnel optimization, and uplift modeling. They operate on Snowflake with dbt-core, Airflow orchestration, and a semantic layer built on cortex.

They have tested multiple AI analytics vendors. The ones that failed tried to replace what the team already did well. The ones that stuck automated the work the team did not want to do. The lesson applies broadly: mature data teams adopt tools that remove noise, not tools that compete with internal talent.

To land Opendoor, Sundial needs to deliver two things:

A. Slack bot for self-service analytics

Data scientists at companies like Opendoor spend roughly 30% of their time answering ad-hoc questions in Slack. "Why did conversion drop in Phoenix?" "What changed in the spread?" "Can you pull the War Room numbers?" These questions interrupt deep work. They pull senior people away from experimentation and modeling.

Sundial's Slack bot intercepts these questions. It runs the analysis in the background, cites the sources, shows the reasoning path, and delivers the answer where the business already works. The data team audits the answers. Over time, the bot handles Tier 1 and Tier 2 questions on its own.

This maps directly to Chandra and Julie's own framework. In their Substack ("Stop Asking Your Data Questions"), they describe three tiers of data push: alerts for emergencies, briefings for morning context,

and ambient discovery for insights nobody thought to ask about. The Slack bot is the entry point for all three.

From Pull to Push: Sundial's Wedge




Figure 1. From pull to push: Sundial's wedge.

B. Semantic model generation via AI-native data catalog

Raw text-to-SQL fails. Chandra wrote about this ("What It Actually Takes to Trust AI With Your Data"): when you point an AI at a warehouse and ask "why did retention drop?", you hand the question to someone who has never worked at your company, does not know which of the 14 retention dashboards uses the canonical definition, and cannot distinguish between a real change and a metric redefinition.

The fix is not a smarter model. It is better context. Canonical metric definitions, business change logs, analytical playbooks, and continuous accuracy measurement. This is the semantic layer.

Data engineers hate writing and maintaining the YAML files that power this layer. Sundial automates semantic model generation directly from the warehouse metadata. It reads the schema, infers business logic, proposes definitions, and lets the data team approve or edit. Once the semantic layer lives in Sundial, every downstream feature gets more accurate. And churn drops to near zero because ripping out the semantic layer means ripping out the foundation.

Chandra's own data: when his team applied proper context to a test set of 50 real user questions, accuracy jumped from about 80% to 98%. Same model. Same warehouse. Same data. The difference was institutional know-how encoded as context.

2. Why the mid-market, and why the market choice is also a product choice

The mid-market is not a compromise between startups and enterprise. It is a distinct category with distinct economics. More importantly, picking this market tells Sundial exactly what to build and what to defer.

The mid-market profile

Attribute	Mid-market reality
Revenue range	\$100M-\$1B
Data team size	10-30 people (DS, AE, DE)
Stack	Modern warehouse (Snowflake/BigQuery), dbt, Airflow
Talent quality	Strong. They hire from FAANG and top programs
Pain point	Time waste, not skill gaps
Buying motion	VP Data / CDO owns budget, 60-90 day sales cycle
Contract value	\$100K-\$500K ARR


Why mid-market beats the alternatives

Startups (\$10M-\$100M revenue, 0-3 data people): The C-Suite needs answers but has nobody to build the infrastructure. Sundial can serve them, but only through a Forward Deployed Engineer model where a Sundial employee sets up the tool, cleans the data, and acts as a fractional analyst. This wins logos fast, but it looks like a consulting business: services-heavy, lower gross margins (40-60%), and higher churn.

Global 2000 (\$1B+ revenue, 50-200+ data people): Sales cycles last 6-12 months. Security reviews take longer than implementation. Worth pursuing eventually, but Sundial does not yet have platform success team to operate at that scale.

The mid-market sweet spot: These companies have budget and urgency. Their data infrastructure is modern enough that Sundial integrates in days, not months. Their teams are strong enough to evaluate the product on merit. And their pain is acute: they are drowning in ad-hoc requests, manual reporting, and semantic layer maintenance.

The Leverage Math



10 Data Scientists × \$250K avg comp × 30% time saved
= **\$750K** annual value recovered

Figure 2. The leverage math behind the mid-market ROI case.

How the market choice simplifies the product roadmap

When a company tries to serve multiple segments simultaneously, the product roadmap fractures. Every segment pulls the product in a different direction. Mid-market companies already use Snowflake and dbt. They need Sundial to sit on top of those tools and add the intelligence layer.

Build now	Defer	Stop
Slack bot with transparent reasoning and admin-managed context	Built-in ETL or data pipeline orchestration	Features that compete with what the internal data team already does well
Semantic model generation from existing dbt models and warehouse metadata	Replacing dbt	Onboarding that asks buyers to change their existing warehouse or transformation layer
Reusable playbooks for funnel reviews, metric movements, and experiment analysis	A full data catalog replacement	Heavy migration-led workflows
Scheduled briefings that lead with what changed	Infrastructure migration flows	Segment-expanding features that dilute the wedge

Roadmap implication: A startup-focused product must build everything from scratch because the customer has nothing. A mid-market product plugs into what exists and adds a thin, high-value layer on top. That layer is faster to build, faster to deploy, and creates structural lock-in through the semantic model.

3. The competitive market map

We map the landscape across two dimensions: how deeply AI is integrated into a company's analytics workflow (AI native-ness) and how mature their data infrastructure is (data maturity). Every company listed below has been verified against its current public positioning as of May 2026.

Where Sundial Wins




Figure 3. Where Sundial wins in the competitive landscape.

Mid-market quadrant: Sundial's target

Delphina - AI-native data platform with a context layer, Slack-native analytics, data agents, scheduled analyses, proactive alerts, and shareable data apps. Their Substack case study shows 35% weekly active users across the company, 5x more questions answered daily, and 95%+ accuracy.

Omni - Governed analytical harness with AI grounded in a semantic layer. Warehouse-first architecture, dbt integration, embedded analytics. Omni frames the moat as the analytical harness, not raw model access. This is Sundial's most direct competitor in the mid-market especially given founder pedigree (e.g. Looker CPO).

Summation - Decision-grade AI platform producing verified, exec-ready reports. Founded by Ian Wong, Opendoor co-founder and CTO. Every figure traces to source data. Reports become repeatable playbooks for WBR and QBR cycles.

PostHog - Open-source product analytics platform moving from reactive dashboards to proactive AI agents. Relevant because it represents the product-analytics-native push into the same proactive insights territory Sundial targets.

Other quadrants

Quadrant	Representative player	Implication for Sundial
Startup: low data maturity, high AI	DatOst (YC): "your first AI data hire"	Useful logo source, but services-heavy and lower-margin unless tightly productized. They are early and my bet is such companies will pivot to AI services.
Global 2000: high data maturity, enterprise scale	TextQL: Ana for petabyte-scale enterprise environments	Eventually attractive, but demands long sales cycles, advanced compliance, and heavier CS
Frontier: high data maturity, high AI	Custom in-house systems such as Ramp Research	Design-partner learning source, but unlikely near-term paying customers
Legacy: low data maturity, low AI	Traditional BI and consultancies	Avoid

4. Sequencing based on target persona

If Sundial targets the mid-market (recommended)

Quarter	Primary move	Execution detail
Q1	Land the beachhead	Ship Slack bot; deploy semantic model generation that reads dbt and warehouse metadata; close 3-5 CDO/VP Data design partners; keep onboarding under 48 hours
Q2	Expand the wedge	Add reusable playbooks for recurring questions; build scheduled briefings; publish case studies showing time saved and accuracy gains
Q3	Lock in and scale	Launch ambient discovery; integrate with in-house data catalog; creates lock-in. Move to \$100K-\$500K ARR contracts

If Sundial targets startups instead

- Q1: Hire 5-10 Forward Deployed Engineers; build out-of-the-box metric templates for ARR, CAC, LTV, and churn; target Series A/B companies with no data team.
- Q2: Standardize the FDE playbook; reduce deployment time to under 72 hours; accept that gross margins will be 40-60% and churn will be higher than mid-market.
- Q3: Use startup logos as social proof to move upmarket.

The procurement pitch

Mid-market procurement teams approve tools that show clear ROI against headcount cost. A data scientist costs \$250K/year fully loaded. If they spend 30% of their time on ad-hoc Slack requests and manual reporting, Sundial recovers \$75K in value per person per year. For a 15-person team, that is \$1.1M in recovered capacity against a \$200K-\$400K Sundial contract. The payback period is under 90 days.

Appendix: alignment with Sundial's own thesis

The arguments in this document align with the philosophy Chandra and Julie have published in Opinionated Intelligence:

Published thesis	How it maps to the blueprint
"What It Actually Takes to Trust AI With Your Data"	The problem is not intelligence, it is context. AI needs canonical metrics, business change logs, analytical playbooks, and continuous accuracy measurement.
"Why AI Analysis Gives You The Wrong Specific Answers"	AI defaults to the most statistically common pattern when context is thin. Sundial's admin-managed context and metric definitions provide the orthogonal context that narrows interpretation to a single correct story.
"Stop Asking Your Data Questions"	The future of data is push: alerts, briefings, discovery. The Slack bot is the entry point, playbooks and briefings are the expansion, and ambient discovery is the long-term moat.

The mid-market is where this thesis has the highest return. These companies have the data maturity to make push work, but they have not built the push infrastructure themselves. They already have the warehouse, the transformation layer, and the talent. What they lack is the automation layer that turns reactive Slack questions into proactive, trusted, recurring intelligence. Sundial fills that gap, and by choosing this market, the product roadmap becomes clear: build on top of what exists, do not replace it.